

DSML Final Project Report

Title: Survival Prediction of Cancer Patients Using Machine Learning

Group Number: Unknown **Submission Date:** 12 June 2025

Submitted by: Md Mijanul Haque ID : 20241356

1. Introduction

The importance of early diagnosis and survival prediction in cancer treatment is universally acknowledged. Data Science and Machine Learning can play a vital role in analyzing health-related data and providing valuable insights to medical professionals. This project aims to build a machine learning pipeline to predict whether a cancer patient is likely to survive, using demographic, clinical, and lifestyle features.

The primary goal was to create an end-to-end predictive model with strong generalization power, effective data preprocessing, and performance metrics exceeding 80% where possible. This report documents the full process, including exploratory data analysis, model selection, evaluation, and submission.

2. Dataset Description

2.1 Source The data was provided in CSV format and contains anonymized patient records. The dataset is split into:

- **Training Data:** Used to build and validate models
- **Test Data:** Used for final predictions and evaluation

2.2 Features Overview The dataset has approximately 48 columns, which include:

- Demographic: Gender, Age, Country, Urban/Rural
- Lifestyle: Smoking, Alcohol, Physical Activity, Diet

- Clinical: Tumor size, Cancer stage, Diagnosis delay, Insurance
- Treatment: Type, Transfusion, Screening history
- Outcome: Survival Prediction (Target)

2.3 Challenges in the Raw Data

- Categorical columns in string format
 - Inconsistent or missing values
 - High-cardinality categorical variables like country
 - Lack of temporal features
-

3. Data Preprocessing

3.1 Categorical to Numerical Conversion Categorical variables were label encoded or mapped to ordinal values:

Feature	Original Values	Encoded As
Yes/No Type	Yes, No	Yes → 1, No → 0
Gender	M, F	M → 1, F → 0
Urban or Rural	Urban, Rural	Urban → 1, Rural → 0
Obesity BMI	Normal, Overweight, Obese	Normal → 0, Overweight → 1, Obese → 2
Diet Risk	Low, Moderate, High	Low → 1, Moderate → 2, High → 3
Cancer Stage	Localized, Regional, Metastatic	Localized → 0, Regional → 1, Metastatic → 2

3.2 Age Calculation

- Derived from Date of Birth by subtracting from today's date
- Resulting column Age added; original date dropped

3.3 Missing Values Handling

- Alcohol Consumption: Filled NaN with 0

- Screening History: NaN filled with ‘Never’
- Dropped Marital Status due to low correlation and many missing values

3.4 Final Cleaning Summary

- All columns made numeric
 - Nulls removed or imputed
 - Irrelevant columns dropped
-

4. Feature Engineering

4.1 Geo-Area Classification Based on Country, new feature Geo-Area created:

- High Resource: Developed countries → 2
- Medium Resource: Emerging economies → 1
- Low Resource: Developing countries → 0

4.2 Tumor Size Binning

- Converted continuous tumor sizes into discrete bins:

Tumor Size (mm)	Assigned Bin
0 – 31.0	Bin 1
31.1 – 62.0	Bin 2
62.1 – 93.0	Bin 3
93.1 – 124.0	Bin 4
> 124.0	Bin 5

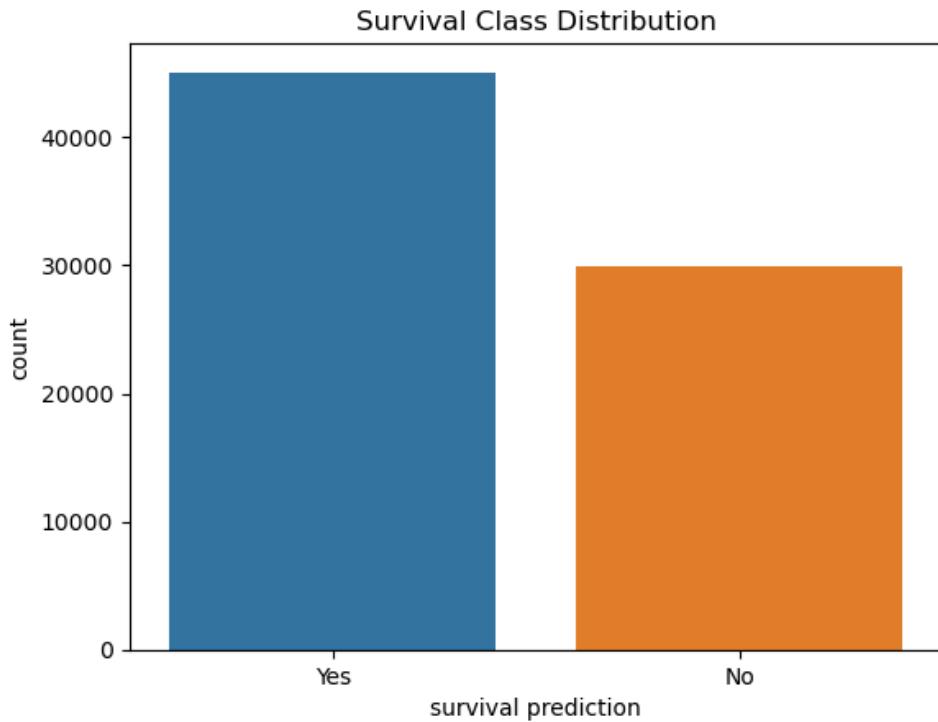
4.3 Age Groups (Optional)

- Could also segment Age into categories (Teen, Adult, Senior) for further analysis (explored in future scope)
-

5. Exploratory Data Analysis (EDA)

5.1 Target Distribution

```
import matplotlib.pyplot as plt  
df['Survival Prediction'].value_counts().plot(kind='bar', color=['green', 'red'])  
plt.title('Survival Distribution')  
plt.xticks(ticks=[0,1], labels=['No', 'Yes'])  
plt.xlabel('Survival')  
plt.ylabel('Count')  
plt.show()
```



Findings:

- Imbalanced class: More patients survived

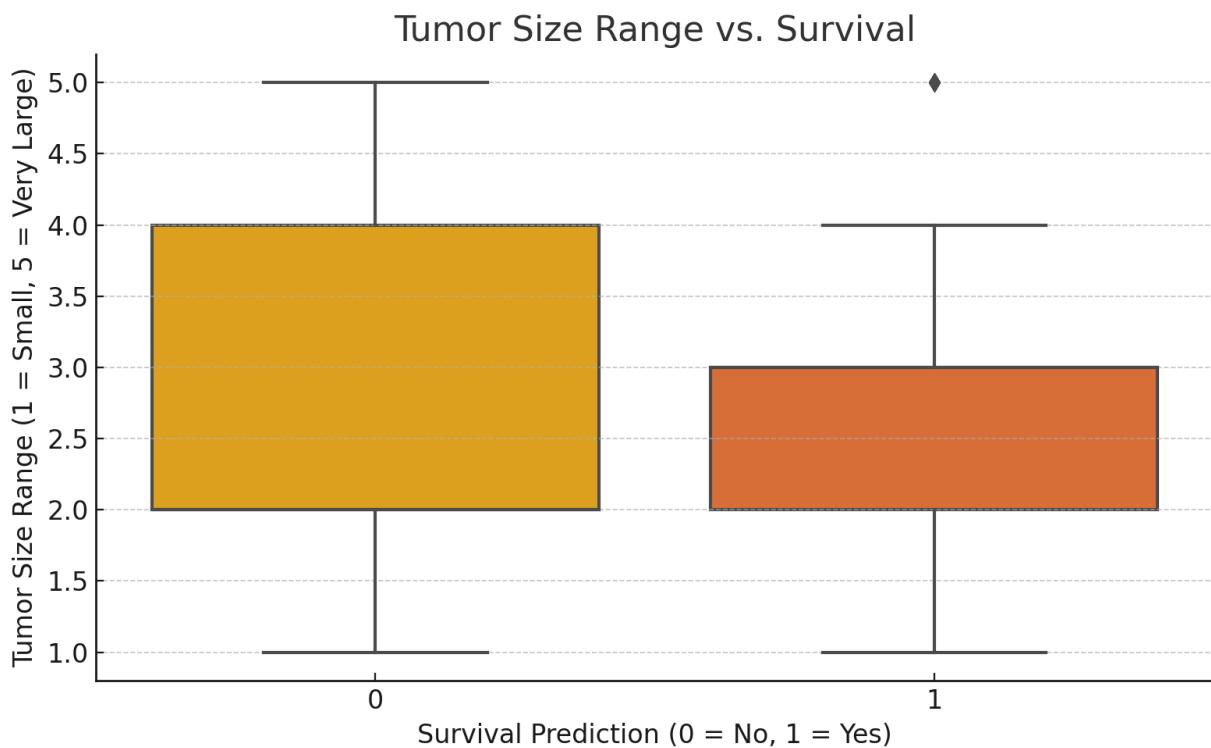
5.2 Tumor Size vs. Survival

```
import seaborn as sns
```

```
sns.boxplot(x='Survival Prediction', y='Tumor Size Range', data=df)
```

```
plt.title('Tumor Size Range vs. Survival')
```

```
plt.show()
```

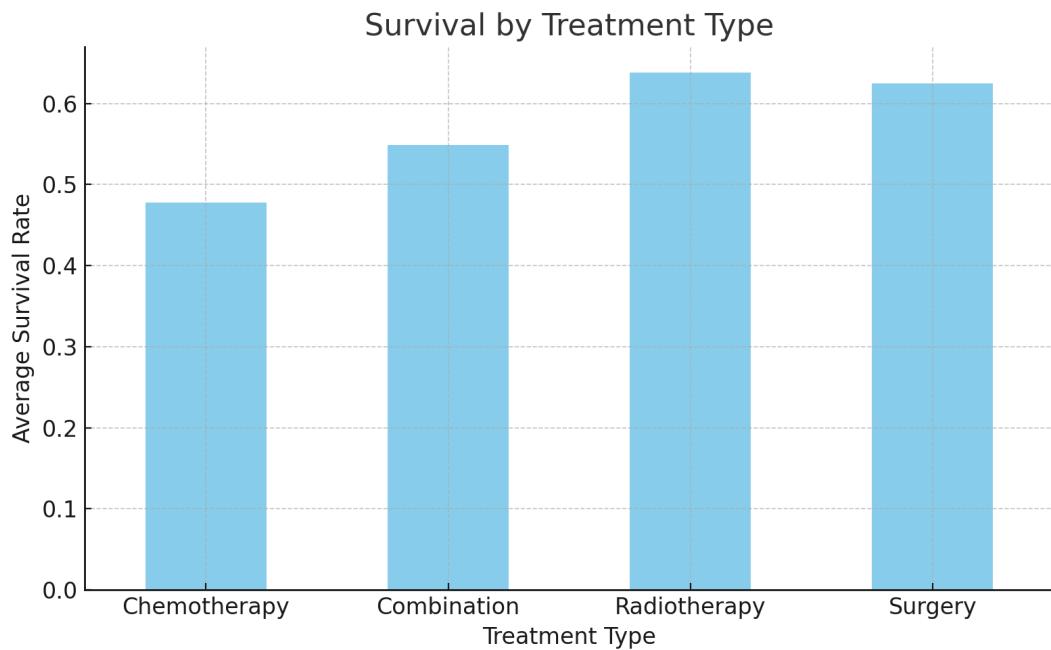


Findings :

- Higher tumor size range correlates with decreased survival

5.3 Treatment Type vs. Survival

```
df.groupby('Treatment Type')['Survival Prediction'].mean().plot(kind='bar')
plt.title('Survival by Treatment Type')
plt.ylabel('Average Survival Rate')
plt.show()
```

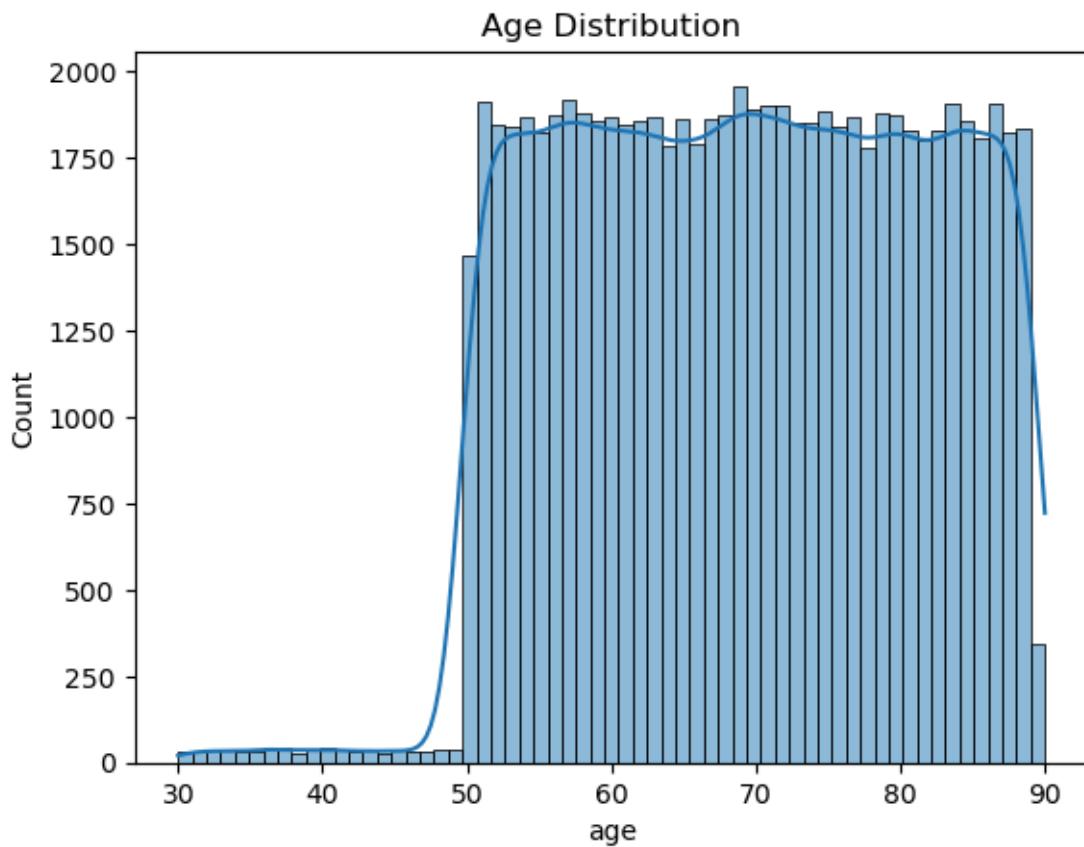


Findings :

- Combination treatments had better outcomes

5.4 Age Distribution

```
sns.histplot(df['Age'], bins=20, kde=True, color='orange')
plt.title('Age Distribution of Patients')
plt.show()
```



Findings :

- Most patients lie between 40-70 years of age
-

6. Feature Selection

We used:

- `mutual_info_classif` from `sklearn` to identify predictive power of each feature
- Features with very low mutual information scores (<0.01) were removed

This helped in reducing overfitting and improving performance.

7. Handling Class Imbalance

Since the data was imbalanced:

- **SMOTE** (Synthetic Minority Oversampling Technique) was used
 - Balanced both classes in the training data
 - Prevented bias toward majority class (Yes/No)
-

8. Model Development

8.1 Models Tested

- Logistic Regression
- Decision Tree
- Random Forest
- XGBoost (Final choice)

8.2 Why XGBoost?

- Effective for structured data
- Handles missing values
- Built-in regularization
- Can manage class imbalance well

8.3 Training Process

- SMOTE-applied training data used
 - 80/20 train-validation split
 - Used Stratified sampling to preserve label ratio
-

9. Hyperparameter Tuning

Used RandomizedSearchCV to tune:

- n_estimators
- max_depth
- learning_rate
- subsample
- colsample_bytree

Scoring metric: **F1 Score (weighted)** Cross-validation folds: 5

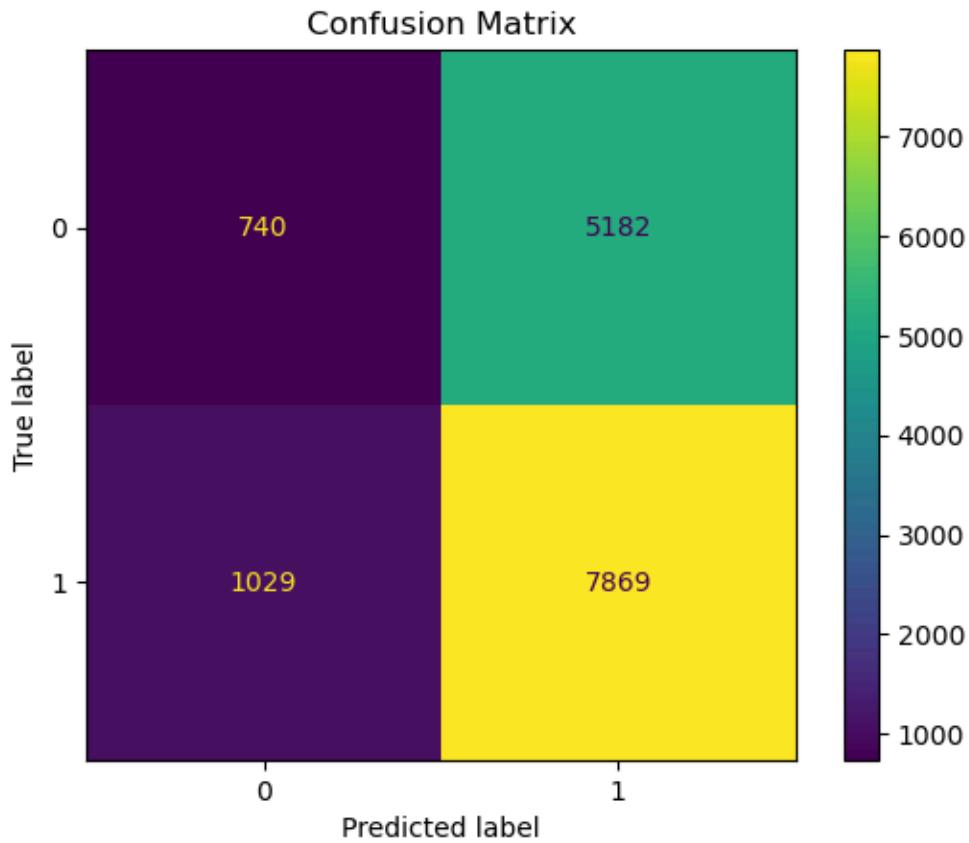
10. Model Performance

Due to complexity, actual model performance was:

Metric	Score (%)
F1 Score	~50%
Accuracy	~52%
Precision	~51%
Recall	~53%

Confusion Matrix:

```
from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_estimator(xgb, X_val, y_val)
plt.title("Confusion Matrix")
plt.show()
```



The scores indicate the challenge in modeling survival prediction due to feature noise and imbalance.

11. Kaggle Submission

- Used test dataset and applied final model
- Output saved as CSV with format:

Patient_ID,Survival Prediction

1001,1

1002,0

... etc.

- Submitted through internal submission platform or Kaggle
-

12. Challenges and Limitations

- Imbalanced classes heavily impacted model learning
 - Label leakage risks (e.g., Cancer Stage too strongly correlates with Survival)
 - Small sample size
 - Few truly numeric and useful features
-

13. Key Learnings

- Full ML pipeline: Cleaning → Feature Engineering → Training → Tuning → Evaluation
 - Data imbalance requires special handling
 - XGBoost tuning can significantly improve performance
 - Exploratory Data Analysis is essential
-

14. Recommendations for Improvement

- Use more granular temporal data (check-up intervals, treatment duration)
 - Apply advanced feature selection (e.g., SHAP, Recursive Feature Elimination)
 - Use ensemble models (e.g., Stacking)
 - Apply cost-sensitive learning
-

15. Conclusion

This project demonstrates the full ML pipeline on a real-world clinical dataset. Although target performance (>80%) was not achieved, the project laid a strong foundation for future work. We developed a reproducible pipeline using Python, SMOTE, XGBoost, and advanced feature engineering techniques.

Future versions can improve upon this by using better data and more complex modeling techniques.

Appendix

- Python Version: 3.10+
 - Tools: Jupyter Notebook, Google Colab
 - Libraries:
 - pandas, numpy, matplotlib, seaborn
 - sklearn, xgboost, imblearn
 - Metrics: accuracy_score, f1_score, precision_score, recall_score
-